

CSC 411

Computer Organization (Spring 2024)

Lecture 16: Compiling/interpreting and running programs

Prof. Marco Alvarez, University of Rhode Island

Context

machine code	assembly	C++	Python
<pre>10110100 10110111 00101011 00100100 00010000 10110111 00010000 11110100 00101000 10101010 00101011 00100001 01010010 11010100 11010001 00101010 10010100 00100000 00010000 00000000 00100000 10101001 00010101 00101010 00100100 10010100 01100001 11110101 11101011 00101111 01010010 10000001 11111110 01010100 10000000 00000000 01010000 11101000 11010001 00100100 10010100 00100000 00000000 00000000 00100001 10101000 00010100 00101010 00010010 10000000 00100000 11110001 11101001 00101011 01010010 10000001 11111110 01010000 00000000 00000000 00000000 00000000 00100000 00010100 00000010 00101010 00101010 00100000 10011111</pre>	<pre>.equ STDOUT, 1 .equ SVC_WRITE, 64 .equ SVC_EXIT, 93 .text .global __start .start: stp x29, x30, [sp, -16]! mov x29, #STDOUT ldr x1, =msg mov x2, _13 mov x8, #SVC_WRITE mov x29, sp svc #0 ldr x29, x30, [sp], 16 mov x8, #0 mov x8, #SVC_EXIT svc #0 msg: .ascii "Hello World!\n" .align 4</pre>	<pre>#include <iostream> int main () { std::cout << "Hello World!" << std::endl; }</pre>	<pre>print('Hello World')</pre>

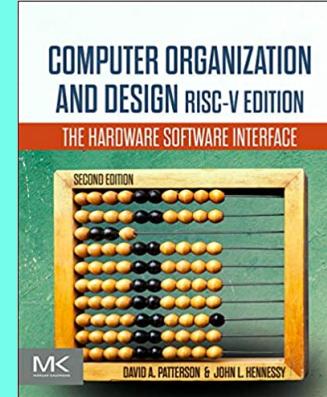
https://www.uvm.edu/~cbcrafier/cs1210/book/02_programming_and_the_python_shell/programming.html

Disclaimer

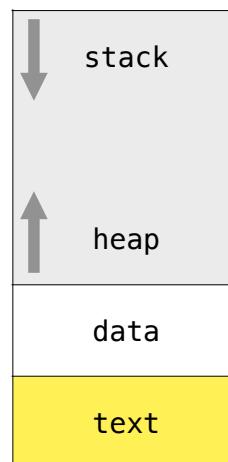
Some figures and slides are adapted from:

Computer Organization and Design (Patterson and Hennessy)

The Hardware/Software Interface



Context



High-level language program
(in C)

Assembly language program (for RISC-V)

```
swap(size_t v[], size_t k)
{
    size_t temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

```

swap:
    slli x6, x11, 3
    add x6, x10, x6
    lw x5, 0(x6)
    lw x7, 4(x6)
    sw x7, 0(x6)
    sw x5, 4(x6)
    jalr x0, 0(x1)

```

Assembler

Compiling vs Interpreting

Compiler

- a program that acts as a **translator**, converting a program from a high-level language into another language
 - most of the time into a low-level language
- because the entire program is translated at once, compilers can perform **optimizations** to make the code more efficient, resulting in faster execution (higher performance)
- e.g. C/C++ compilers

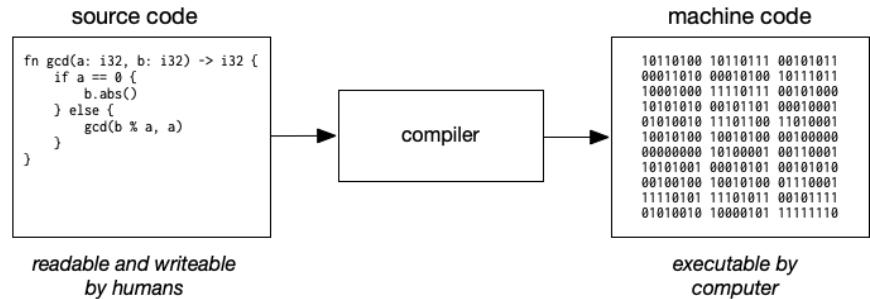
Interpreter

- a program that reads and **translates code on the fly**
 - reads code line by line, translates it into machine code, and executes it
 - any language can be interpreted (including C or assembly)
- preferred when performance is not critical
- e.g. Javascript

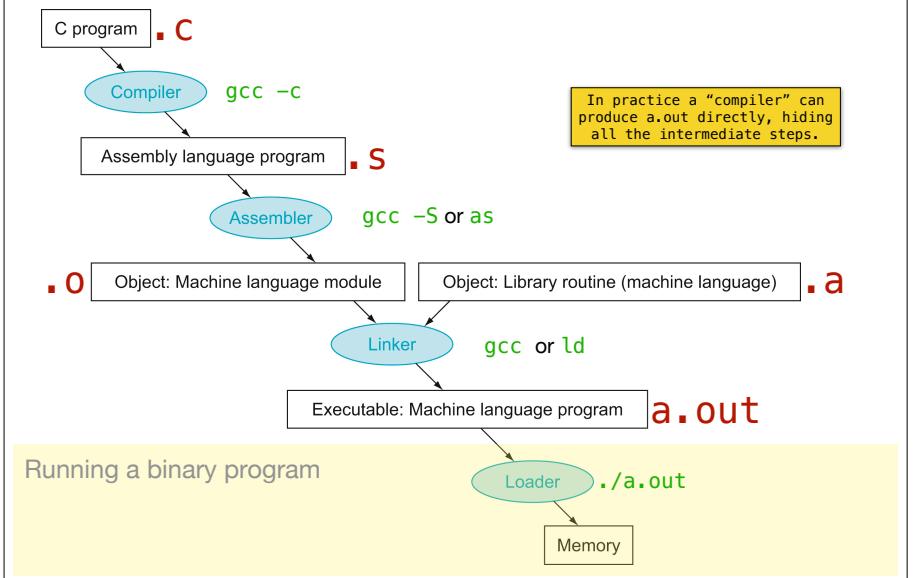
Feature	Compiling	Interpreting
Translation Time	Entire program translated before execution	Code translated and executed line by line
Output	Machine code (executable file)	No direct output; relies on interpreter for execution
Execution Speed	Generally faster due to pre-optimization	Generally slower due to on-the-fly translation
Development Speed	Slower due to separate compilation step	Faster; code can be run and tested immediately
Portability	Limited portability (requires recompilation for different architectures)	Higher portability; interpreted code can often run on different systems
Optimizations	Extensive optimizations possible during compilation	Limited optimizations during interpretation
Error Checking	Can perform more comprehensive static analysis	Limited error checking during interpretation
Code Distribution	Compiled machine code (binary files)	Source code in original language

Compiling

Compiling programs (simplified)



Compiling/linking/running C programs



C Compiler

- Translates C language into assembly
 - symbolic form of machine language
 - in the 70s many operating systems were written in assembly language
 - Optimizes generated code
 - optimization levels are specified using flags like `-O1`, `-O2`, and `-O3`, which progressively increase the level of optimization performed by the compiler
 - other optimization flags include `-O0`, `-Os`, `-Ofast`, `-Og`, `-Oz`

<https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

C Compiler

Optimization name	Explanation	gcc level
<i>High level</i>	<i>At or near the source level; processor independent</i>	
Procedure integration	Replace procedure call by procedure body	03
<i>Local</i>	<i>Within straight-line code</i>	
Common subexpression elimination	Replace two instances of the same computation by single copy	01
Constant propagation	Replace all instances of a variable that is assigned a constant with the constant	01
Stack height reduction	Rearrange expression tree to minimize resources needed for expression evaluation	01
<i>Global</i>	<i>Across a branch</i>	
Global common subexpression elimination	Same as local, but this version crosses branches	02
Copy propagation	Replace all instances of a variable A that has been assigned X (i.e., $A = X$) with X	02
Code motion	Remove code from a loop that computes the same value each iteration of the loop	02
Induction variable elimination	Simplify/eliminate array addressing calculations within loops	02
<i>Processor dependent</i>	<i>Depends on processor knowledge</i>	
Strength reduction	Many examples; replace multiply by a constant with shifts	01
Pipeline scheduling	Reorder instructions to improve pipeline performance	01
Branch offset optimization	Choose the shortest branch displacement that reaches target	01

Assembler

- Converts pseudo-instructions into actual instructions
 - Converts actual assembly instructions into **machine code**
 - must determine memory addresses and offsets corresponding to all labels (requires two passes over program)
 - Generates an **object file**
 - combination of machine instructions, data, and information needed to place instructions properly in memory

An object file

- Intermediate file produced by a compiler, typically consisting of the following:
 - object file header:** describes contents (size/position) of other pieces of the object
 - text segment:** translated instructions (machine code)
 - static data segment:** data allocated for the life of the program
 - relocation information:** for instructions/data that depend on absolute addresses (updated by the linker later)
 - symbol table:** list of symbols (functions, labels, etc.) along with their respective addresses or values that can be used by other objects
 - debugging information:** metadata for associating machine code with original source code

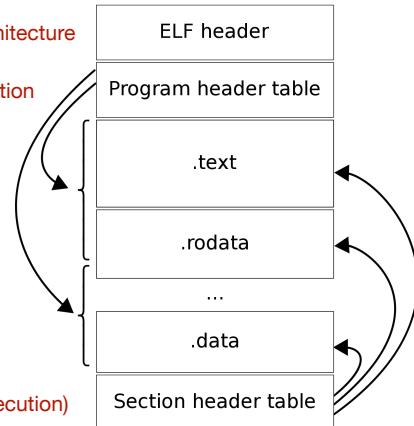
ELF files

Executable and Linkable Format

Identify the ELF type and specify the architecture

Execution information

Technical details for linking (ignored for execution)

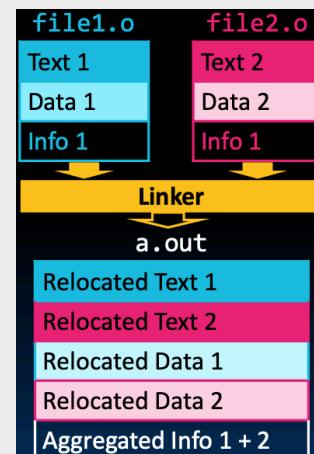


https://en.wikipedia.org/wiki/Executable_and_Linkable_Format

Linker

- Merges all object files, producing an executable
 - merge code and data segments from all objects
 - resolve all undefined labels (replace them with final addresses)
 - patch location-dependent and external references
- Could leave location dependencies to be fixed by a relocating loader
 - with virtual memory, no need to do this
 - program can be loaded into absolute location in virtual memory space

Linker



From Berkeley's CS61C lectures

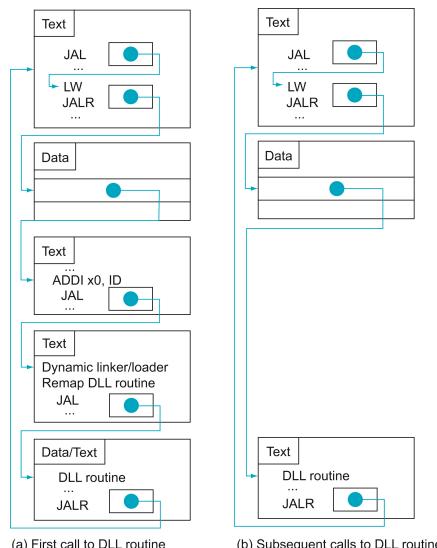
Loader (OS task)

- (1) reads header to determine text/data segment sizes
- (2) creates address space large enough for text/data
- (3) copies text/data from executable into memory
- (4) copies the command line arguments (if any) onto the stack
- (5) initializes registers (including sp, fp, gp)
- (6) jump to startup routine
 - copies arguments to x10 and x11 and calls main
 - when main returns, terminate program with an exit system call

Lazy linkage

Dynamically linked library via lazy procedure linkage:

- (a) Steps for the first time a call is made to the DLL routine.
- (b) The steps to find the routine, remap it, and link it are skipped on subsequent calls.



Dynamically linked libraries

• Statically linked libraries

- library procedures become part of the executable code
 - if new versions are released code must be recompiled otherwise executable keeps using the old versions
- loads procedures in the library even if those calls are not executed

• Dynamically linked libraries (DLLs)

- only link/load library procedure when it is called
- automatically picks up new library versions
- requires procedure code to be relocatable

Example

sum_int.h

```
#ifndef __SUM__
#define __SUM__

int sum(int a, int b);
#endif
```

sum_int.c

```
extern int baseline;

int sum(int a, int b) {
    if (a < b)
        return (a + b) - baseline;
    else
        return a + b;
}
```

prog.c

```
#include "sum_int.h"

int baseline = 10;

int main() {
    int a = 10;
    int b = 20;

    int c = sum(a, b);

    return 0;
}
```

sum_int.o

```
xxd -q4 -c 32 sum int.o
```

```
objdump -d sum int.o
```

```

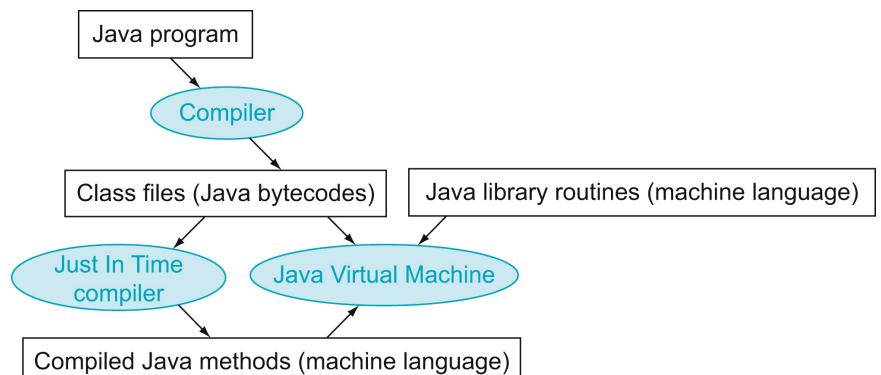
00000000 <sum>:
 8: 0fe10131 addisp,sp,-32
 4: 0081e2e23 sw $0,(sp)
 8: 02010413 addis$0,sp,32
 3: fea42623 sw $0,-20($0)
10: fea42703 lw $0,24($0)
14: fea42703 lw $0,-20($0)
18: fe842783 lw $0,-24($0)
1c: 02f75063 bge a4,$5,3c <L2>
20: fea42703 lw $0,24($0)
24: fea42783 lw $0,-24($0)
28: 00817873 add a4,a4,a5
32: 0080007d lui a5,$0x8
30: 00807a83 lw a5,$0(a5) # 0 <sum>
34: 46178053 sub a5,a4,a5
38: 0000006f j 48 <L3>

0000003c <_L2>:
 3: fea42703 lw a4,-20($0)
40: fe842783 lw $0,-24($0)
44: 00817873 add a5,a4,a5

```

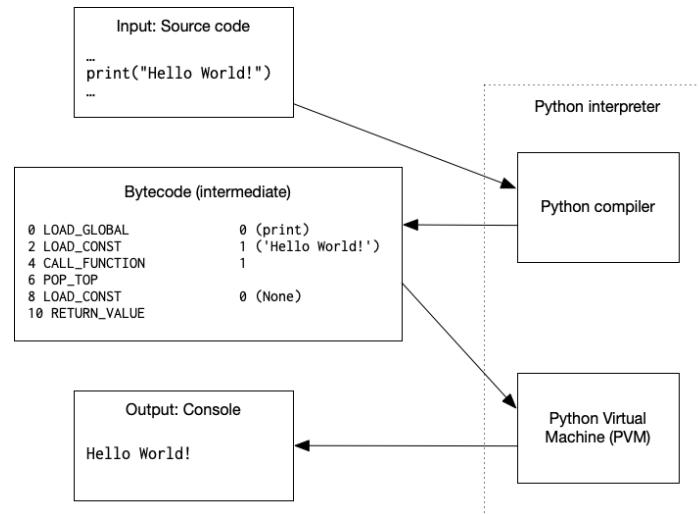
```
objdump -d a.out
```

Java applications



Modern interpreters

Python interpreter



Python interpreter

```
mylist = [1, 2, 3, 'hello', 'world']

myslice = mylist[3:]

print(myslice)
```

python -m dis <file>

```

1  0 BUILD_LIST      0 ((1, 2, 3, 'hello', 'world'))
2  2 LOAD_CONST      0 (1, 2, 3, 'hello', 'world')
4  4 LIST_EXTEND    1
6  6 STORE_NAME      0 (mylist)

2  8 LOAD_NAME        0 (mylist)
10 10 LOAD_CONST     1 (3)
12 12 LOAD_CONST     2 (None)
14 14 BINARY_SUBSCR 2
16 16 STORE_NAME      1 (myslice)

3  20 LOAD_NAME       2 (print)
22 22 LOAD_NAME       1 (myslice)
24 24 CALL_FUNCTION   1
26 26 POP_TOP
28 28 LOAD_CONST     2 (None)
30 30 RETURN_VALUE
```

Impacts on performance

Sort example in C

```
void swap(int v[], size_t k) {
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}

void sort (int v[], size_t n) {
    size_t i, j;
    for (i = 0; i < n; i += 1) {
        for (j = i - 1; j >= 0 && v[j] > v[j + 1] ; j -= 1) {
            swap(v, j);
        }
    }
}
```

```

swap:
    slli x6, x11, 2 # reg x6 = k * 4
    add x6, x10, x6 # reg x6 = v + (k * 4)
    lw x5, 0(x6) # reg x5 (temp) = v[k]
    lw x7, 4(x6) # reg x7 = v[k + 1]
    sw x7, 0(x6) # v[k] = reg x7
    sw x5, 4(x6) # v[k+1] = reg x5 (temp)
    jalr x0, 0(x1) # return to caller

sort:
    forl1st: bge x19, x22, exit1
    addi x20, x19, -1 # j = i - 1
    for2tst: blt x20, x0, exit2
    addi x19, x20, 1 # j = j + 1
    slli x5, x20, 2 # x5 = j * 4
    add x5, x21, x5 # x5 = v + (j * 4)
    lw x6, 0(x5) # x6 = v[j]
    lw x7, 4(x5) # x7 = v[j + 1]
    swap(x6, x7)
    ble x6, x7, exit2
    addi x10, x21, 0 # first swap parameter is v
    addi x11, x20, 0 # second swap parameter is j
    jal x1, swap # call swap
    for2tst:
    addi x20, x19, -1 # j for2tst
    jal x0, for2tst # go to for2tst
    forl1st:
    addi x19, x19, 1 # i += 1
    jal x0, forl1st # go to forl1st
```

Saving registers	
sort:	addi sp, sp, -20 # make room on stack for 5 registers sw x1, 16(sp) # save return address on stack sw x22, 12(sp) # save x22 on stack sw x21, 8(sp) # save x21 on stack sw x20, 4(sp) # save x20 on stack sw x19, 0(sp) # save x19 on stack
Procedure body	
Move parameters	addi x21, x10, 0 # copy parameter x10 into x21 addi x22, x11, 0 # copy parameter x11 into x22
Outer loop	forl1st: bge x19, x22, exit1 # go to exit1 if i >= n
Inner loop	addi x20, x19, -1 # j = i - 1 for2tst: blt x20, x0, exit2 # go to exit2 if j < 0 slli x5, x20, 2 # x5 = j * 4 add x5, x21, x5 # x5 = v + (j * 4) lw x6, 0(x5) # x6 = v[j] lw x7, 4(x5) # x7 = v[j + 1] ble x6, x7, exit2 # go to exit2 if x6 < x7
Pass parameters and call	addi x10, x21, 0 # first swap parameter is v addi x11, x20, 0 # second swap parameter is j jal x1, swap # call swap
Inner loop	addi x20, x19, -1 # j for2tst jal x0, for2tst # go to for2tst
Outer loop	exit2: addi x19, x19, 1 # i += 1 jal x0, forl1st # go to forl1st
Restoring registers	
exit1:	lw x19, 0(sp) # restore x19 from stack lw x20, 4(sp) # restore x20 from stack lw x21, 8(sp) # restore x21 from stack lw x22, 12(sp) # restore x22 from stack lw x1, 16(sp) # restore return address from stack addi sp, sp, 20 # restore stack pointer
Procedure return	
	jalr x0, 0(x1) # return to calling routine

Applying compiler optimizations

gcc optimization	Relative performance	Clock cycles (millions)	Instruction count (millions)	CPI
None	1.00	158,615	114,938	1.38
O1 (medium)	2.37	66,990	37,470	1.79
O2 (full)	2.38	66,521	39,993	1.66
O3 (procedure integration)	2.41	65,747	44,993	1.46

The programs sorted 100,000 32-bit words with the array initialized to random values. These programs were run on a Pentium 4 with a clock rate of 3.06 GHz and a 533 MHz system bus with 2 GB of PC2100 DDR SDRAM. It used Linux version 2.4.20.

Optimization example

Comparing sorting algorithms

- Compiler optimizations are sensitive to the algorithm
- Java/JIT significantly faster than JVM interpreted

Language	Execution method	Optimization	Bubble Sort relative performance	Quicksort relative performance	Speedup Quicksort vs. Bubble Sort
C	Compiler	None	1.00	1.00	2468
	Compiler	O1	2.37	1.50	1562
	Compiler	O2	2.38	1.50	1555
	Compiler	O3	2.41	1.91	1955
	Interpreter	-	0.12	0.05	1050
Java	JIT compiler	-	2.13	0.29	338

Nothing can fix slow algorithms

Arrays vs pointers

- Array indexing typically involves
 - multiplying index by element size
 - adding to array base address
- Pointers correspond directly to memory addresses
 - can avoid indexing complexity

Example: clearing an array

```
void clear1(int array[], int size) {    void clear2(int *array, int size) {  
    int i;        int *p = array;  
    for (i = 0 ; i < size ; i += 1)    int *end = p + size;  
        array[i] = 0;            for ( ; p < end ; p++)  
    }                                *p = 0;  
}
```

```
l1: li x5, 0      # i = 0  
slli x6, x5, 2    # x6 = i * 4  
add x7, x10, x6 # x7 = &array[i]  
sw x0, 0(x7)     # array[i] = 0  
addi x5, x5, 1   # i = i + 1  
blt x5, x11, l1 # if (i<size) go to l1  
  
l2: add x5, x0, x10 # p = addr of array[0]  
slli x6, x11, 2 # x6 = size * 4  
add x7, x10, x6 # x7 = addr of array[size]  
sw x0, 0(x5)     # *p = 0  
addi x5, x5, 4   # p = p + 4  
bltu x5, x7, l2 # if (p<end) go to l2
```

(*) assembly code above does not consider the case of size = 0

compilers can do this by applying optimizations: **strength reduction** (shift instead of multiply) and **induction variable elimination** (eliminate address calculation inside the loop)

Fallacies

- Powerful instructions (x86) ⇒ higher performance
 - complex instructions are hard to implement
 - may slow down all instructions, including simple ones
 - compilers are good at making fast code from simple instructions
- Write assembly code for higher performance
 - modern compilers are better (than most humans) at dealing with modern processors
 - more lines of code ⇒ more errors and less productivity

Fallacies

- Backward compatibility ⇒ instruction set “doesn’t change”
 - x86 keeps adding more instructions (difficult to build compatible processors)

Growth of the x86 instruction set over time

